

Zapier Monthly Active User Analysis

January 20, 2020

0.0.1 Importing modules and packages

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
import psycopg2
import prophet
from fbprophet import Prophet
from fbprophet.plot import plot_plotly
import plotly.offline as py
import os
warnings.filterwarnings('ignore')
%matplotlib inline
py.init_notebook_mode()
color = sns.color_palette()
pd.set_option('display.float_format', lambda x: '%.3f' % x)
sns.set(rc={'figure.figsize':(8,5)})
```

0.0.2 Connecting to the Zapier database

- A view called **active_users_analysis** has been created in the **jchumley** schema using the *tasks_used_da* table.
- We'll use the *psycopg* odbc module to connect to the Redshift cluster. Pandas will be used extensively in conjunction with Postgres/Redshift queries.
- The user, password, host, and database parameters are loaded from a .bashrc file.

```
[2]: port='5439'
user=os.environ['ZAP_USER']
password=os.environ['ZAP_PASS']
host=os.environ['ZAP_HOST']
db=os.environ['ZAP_DB']
conn = psycopg2.connect(user=user, password=password, host=host, port=port,
↳database=db)
```

0.0.3 Creating dataframe

- Let's create our first dataframe using pandas read_sql method.

```
[3]: query = """
      SELECT * FROM
      jchumley.active_users_analysis
      """
      df = pd.read_sql(query, conn)
      df.head()
```

```
[3]:      date  user_id  account_id  sum_tasks_used  cum_sum_tasks_used  \
0  2017-01-01         1           1             48                48
1  2017-01-02         1           1             65               113
2  2017-01-03         1           1             71               184
3  2017-01-04         1           1             64               304
4  2017-01-05         1           1             74               406

      num_days_since_last_active  active  churned  month_num  month  dow_num  \
0                             0    True   False     1.000  January         0
1                             1    True   False     1.000  January         1
2                             1    True   False     1.000  January         2
3                             1    True   False     1.000  January         3
4                             1    True   False     1.000  January         4

      dow
0  Sunday
1  Monday
2  Tuesday
3  Wednesday
4  Thursday
```

- What is the shape of this dataset?

```
[4]: df.shape
```

```
[4]: (10547587, 12)
```

- There are ~10.5m rows and 12 columns.
- We'll also compute the ratio of missing values for each column.

```
[5]: df_na = (df.isnull().sum() / len(df)) * 100
      df_na = df_na.drop(df_na[df_na == 0].index).sort_values(ascending=False)
      missing_data = pd.DataFrame({'Missing Ratio' :df_na})
      missing_data.head()
```

```
[5]: Empty DataFrame
      Columns: [Missing Ratio]
```

Index: []

- Great, we don't have any null values!

0.1 Exploratory data analysis

0.1.1 Correlation Matrix

- Let's gain a deeper understanding of the associations between the features in this dataset. We'll use the pandas *corr* method to complete this task.

```
[6]: df.corr()
```

```
[6]:
```

	user_id	account_id	sum_tasks_used	\
user_id	1.000	0.974	-0.022	
account_id	0.974	1.000	-0.019	
sum_tasks_used	-0.022	-0.019	1.000	
cum_sum_tasks_used	0.984	0.958	-0.022	
num_days_since_last_active	0.021	0.020	-0.012	
active	-0.008	-0.008	0.004	
churned	0.007	0.006	-0.003	
month_num	0.108	0.108	0.001	
dow_num	0.002	0.002	0.000	

	cum_sum_tasks_used	num_days_since_last_active	\
user_id	0.984	0.021	
account_id	0.958	0.020	
sum_tasks_used	-0.022	-0.012	
cum_sum_tasks_used	1.000	0.023	
num_days_since_last_active	0.023	1.000	
active	-0.009	-0.760	
churned	0.008	0.519	
month_num	0.096	0.060	
dow_num	0.002	-0.031	

	active	churned	month_num	dow_num
user_id	-0.008	0.007	0.108	0.002
account_id	-0.008	0.006	0.108	0.002
sum_tasks_used	0.004	-0.003	0.001	0.000
cum_sum_tasks_used	-0.009	0.008	0.096	0.002
num_days_since_last_active	-0.760	0.519	0.060	-0.031
active	1.000	-0.873	-0.041	-0.000
churned	-0.873	1.000	0.029	0.000
month_num	-0.041	0.029	1.000	0.016
dow_num	-0.000	0.000	0.016	1.000

- For this stage of analysis, the *user_id* and *account_id* columns will be ignored since they are unique identifiers.

- Notice that there is a strong negative association between the features active and churn which makes sense since they are mutually exclusive.

0.1.2 Active and churned users by month

- Let's take a look at counts of active and churned users.
- We'll run the following aggregate query to view churn and active user counts for each month in the dataset.

```
[7]: query_month = """
SELECT month, month_num,
       count(CASE WHEN churned is TRUE then 1 END) as churned_count,
       count(CASE WHEN active is TRUE then 1 END) as active_count,
       ROUND(100.0*count(CASE WHEN churned is TRUE then 1 END)/count(*),2) as_
→churned_ratio,
       ROUND(100.0*count(CASE WHEN active is TRUE then 1 END)/count(*),2) as_
→active_ratio,
       count(*) as total_user_count
FROM jchumley.active_users_analysis
GROUP BY month, month_num
ORDER BY month_num;
"""
df_month = pd.read_sql(query_month, conn)
```


```
[8]: df_month
```

```
[8]:      month  month_num  churned_count  active_count  churned_ratio  \
0   January      1.000          120      1927467         0.010
1  February      2.000          8256      1881695         0.440
2   March       3.000         12974      2207687         0.580
3   April       4.000         12676      2100295         0.600
4    May       5.000         14207      2282399         0.620
5   June       6.000           964       83651         1.140

      active_ratio  total_user_count
0          99.990          1927587
1          99.560          1889981
2          99.300          2223331
3          99.170          2117882
4          99.070          2303880
5          98.500           84926
```

```
[9]: sns.barplot(x='month', y='active_count', palette="rocket", data=df_month)
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x1a374513d0>
```



- We can see that May had the highest count of active users. June has a very low count, perhaps because of data availability. Let's verify.

```
[10]: query_june = """
SELECT month, COUNT(DISTINCT(date))
FROM jchumley.active_users_analysis
GROUP BY month"""
pd.read_sql(query_june, conn)
```


```
[10]:   month  count
0  January     31
1    June       1
2     May     31
3  April     30
4 February     28
5  March     31
```

- We can see that the dataset includes records from the first day of June 2017 but not other days in that month/year.

0.1.3 Churn by month

```
[11]: sns.barplot(x='month', y='churned_count', palette="rocket", data=df_month)
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x1a38a6a310>
```



- January's churn count would likely increase if records from 2016 were available.
- Potential factors contributing to June's relatively high churn count:
 1. Many users either chose to not renew or cancel their subscriptions on the first day of a given month.
 2. All June records are from the first day of the month.

0.1.4 Day of week analysis

- We'll run a query similar to the one used to compute aggregates by month.

```
[12]: dow_query = """
SELECT dow, dow_num,
       count(CASE WHEN churned is TRUE then 1 END) as churned_count,
       count(CASE WHEN active is TRUE then 1 END) as active_count,
       ROUND(100.0*count(CASE WHEN churned is TRUE then 1 END)/count(*),2) as_
→churned_ratio,
       ROUND(100.0*count(CASE WHEN active is TRUE then 1 END)/count(*),2) as_
→active_ratio,
       count(*) as total_user_count
FROM jchumley.active_users_analysis
GROUP BY dow, dow_num
ORDER BY dow_num;
"""
df_dow = pd.read_sql(dow_query, conn)
```

```
[13]: df_dow
```

```
[13]:
```


	dow	dow_num	churned_count	active_count	churned_ratio	\
0	Sunday	0	3619	1102657	0.330	
1	Monday	1	8207	1657615	0.490	
2	Tuesday	2	8524	1703666	0.500	
3	Wednesday	3	9268	1708853	0.540	
4	Thursday	4	8773	1689305	0.520	
5	Friday	5	6805	1540253	0.440	
6	Saturday	6	4001	1080845	0.370	

	active_ratio	total_user_count
0	99.570	1107429
1	99.370	1668200
2	99.340	1714992
3	99.300	1720901
4	99.320	1700834
5	99.420	1549246
6	99.530	1085985

0.1.5 Counts of active users by day of week

```
[14]: sns.barplot(x='dow', y='active_count', palette="rocket", data=df_dow)
```


```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0x1a39d85550>
```



0.1.6 Counts of churned users by day of week

```
[15]: sns.barplot(x='dow', y='churned_count', palette="rocket", data=df_dow)
```

```
[15]: <matplotlib.axes._subplots.AxesSubplot at 0x1a3a400050>
```



- Since Zapier is a B2B platform, its users are expected to be most active while they are in the office.
- Recall that **churned** and **active** have a very strong negative association (see the correlation matrix above), so it's expected that churn is higher during normal business hours.

0.1.7 Power Users

- Let's see who the top 10 Zapier users based on the sum_tasks_used column.
- We'll query the view total_sum_tasks_used which has the cumulative total of sum of tasks used for each user from the tasks_used_da table.

```
[16]: power_user_query = """SELECT user_id, total_sum_tasks_used,
    ROUND(100.0*total_sum_tasks_used/(SELECT SUM(total_sum_tasks_used) as
    ↳total_sum_tasks_used_all_users
FROM jchumley.total_sum_tasks_used),2) pct_all_sum_tasks_used,
    RANK() OVER(ORDER BY total_sum_tasks_used DESC) as
    ↳user_rank_by_sum_tasks_used
FROM jchumley.total_sum_tasks_used
LIMIT 10"""
```



```
df_power_users = pd.read_sql(power_user_query, conn)
df_power_users
```

```
[16]: user_id  total_sum_tasks_used  pct_all_sum_tasks_used  \
0    541993             15878183             3.120
1     19415             12574205             2.470
2    558906             8282192             1.630
3    645563             6153468             1.210
4    103042             3872036             0.760
5   1830742             3305466             0.650
6    166842             3136333             0.620
7    289586             2659815             0.520
8    617995             2335859             0.460
9    115417             2174002             0.430
```

```
user_rank_by_sum_tasks_used
0             1
1             2
2             3
3             4
4             5
5             6
6             7
7             8
8             9
9            10
```

- Would be interesting to dig deeper into metadata on these users (e.g, age, industry, gender, profession, etc.)

0.2 Timeseries forecasting

- We'll use Prophet to forecast active users and churn into the next 365 days. See <https://facebook.github.io/prophet/> if you're not familiar with Prophet.
- Let's run another aggregate query against the active_users_analysis view, this time we'll group by date.

```
[17]: user_aggregates_query = """
SELECT date as ds,
       count(CASE WHEN churned is TRUE then 1 END) as churned_count,
       count(CASE WHEN active is TRUE then 1 END) as active_count,
       ROUND(100.0*count(CASE WHEN churned is TRUE then 1 END)/count(*),2) as_
→churned_ratio,
       ROUND(100.0*count(CASE WHEN active is TRUE then 1 END)/count(*),2) as_
→active_ratio,
       count(*) as total_count
FROM jchumley.active_users_analysis
GROUP BY date
```

```
ORDER BY date"""
df_ts = pd.read_sql(user_aggregates_query, conn)
df_ts.head()
```

```
[17]:
```

	ds	churned_count	active_count	churned_ratio	active_ratio	\
0	2017-01-01	0	38903	0.000	100.000	
1	2017-01-02	0	53046	0.000	100.000	
2	2017-01-03	0	64350	0.000	100.000	
3	2017-01-04	0	65926	0.000	100.000	
4	2017-01-05	0	66304	0.000	100.000	

	total_count
0	38903
1	53046
2	64350
3	65926
4	66304

0.2.1 Forecast of Active Users

- Let's create a timeseries forecast of counts of active users into June 2018.

```
[18]: df_ts_active = df_ts.copy()
df_ts_active[['ds', 'y']] = df_ts[['ds', 'active_count']]
```

Instantiating and Prophet object and making predictions into June 2018.

```
[19]: m = Prophet()
m.fit(df_ts_active)
active_users_future = m.make_future_dataframe(periods=365)
active_users_future.tail()
```

INFO:fbprophet:Disabling yearly seasonality. Run prophet with yearly_seasonality=True to override this.
 INFO:fbprophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to override this.

```
[19]:
```

	ds
512	2018-05-28
513	2018-05-29
514	2018-05-30
515	2018-05-31
516	2018-06-01

```
[20]: active_users_forecast = m.predict(active_users_future)
active_users_forecast[['ds', 'yhat', 'yhat_lower', 'yhat_upper']].tail()
```

```
[20]:
```

	ds	yhat	yhat_lower	yhat_upper
512	2018-05-28	93785.262	65043.437	123242.307
513	2018-05-29	95813.956	66943.975	125219.928


```

514 2018-05-30 95986.144    66981.395   125528.416
515 2018-05-31 95036.977    65946.426   124151.378
516 2018-06-01 91740.225    63269.911   121109.604


```

Visualizing our active users forecast

```
[21]: fig = m.plot(active_users_forecast)
```



```
[22]: fig2 = m.plot_components(active_users_forecast)
```



- The model is predicting a steady growth of active users into June 2018.
- The weekly plot component shown above aligns with our day of week analysis.

0.2.2 Forecast of Churned Users

Similarly, let's create a timeseries forecast of counts of churned users into June 2018.

```
[23]: df_ts_churned = df_ts.copy()
df_ts_churned[['ds', 'y']] = df_ts[['ds', 'churned_count']]
df_ts_churned.head()
```

```
[23]:
```

	ds	churned_count	active_count	churned_ratio	active_ratio	\
0	2017-01-01	0	38903	0.000	100.000	
1	2017-01-02	0	53046	0.000	100.000	
2	2017-01-03	0	64350	0.000	100.000	
3	2017-01-04	0	65926	0.000	100.000	
4	2017-01-05	0	66304	0.000	100.000	

	total_count	y
0	38903	0
1	53046	0
2	64350	0
3	65926	0
4	66304	0

Instantiating and Prophet object and making predictions

```
[24]: m2=Prophet()  
m2.fit(df_ts_churned)  
churned_future = m2.make_future_dataframe(periods=365)  
churned_future.tail()
```

```
INFO:fbprophet:Disabling yearly seasonality. Run prophet with  
yearly_seasonality=True to override this.  
INFO:fbprophet:Disabling daily seasonality. Run prophet with  
daily_seasonality=True to override this.
```


```
[24]:          ds  
512 2018-05-28  
513 2018-05-29  
514 2018-05-30  
515 2018-05-31  
516 2018-06-01
```

```
[25]: churned_user_forecast = m2.predict(churned_future)  
churned_user_forecast[['ds', 'yhat', 'yhat_lower', 'yhat_upper']].tail()
```


```
[25]:          ds      yhat  yhat_lower  yhat_upper  
512 2018-05-28  999.565    408.261    1629.298  
513 2018-05-29 1011.668    398.676    1647.355  
514 2018-05-30 1043.238    411.981    1669.813  
515 2018-05-31 1018.547    405.901    1648.968  
516 2018-06-01  949.444    352.415    1545.710
```

Visualizing out churned users forecast

```
[26]: fig = m2.plot(churned_user_forecast)
```



[28]: `fig2 = m.plot_components(churned_user_forecast)`



- Similar to the forecast of active users, churn is predicted to have steady growth into June 2018 but also at a higher rate.
- Unlike the active user plot, we can see outliers in churn at the end of each month. This could be due to a large number of users who are registered under monthly billing subscriptions.

0.2.3 Business recommendations

- Identify new data sources related to Zapier user metadata (e.g., registration date, plan type, first usage date, acquisition channel, date of birth, gender, industry, etc.)
- Using additional data source combined with the task_used_da table, develop user metadata features
- Develop a logistic regression model to generate churn probability scores for each user
- Perform cohort analysis to group users by shared attributes
- For users who have churned, send out a survey to understand why they have decided to no longer use the Zapier platform
- Using the logistic regression model, identify high-value power users who are at risks of churn
- Run A/B testing experiments and compare churn rates between the test and control slices
- Reach out to users who are on monthly subscription plans prior to the end of the month via email or another communication channel with the goal of customer retention