Zapier Monthly Active User Analysis

January 20, 2020

0.0.1 Importing modules and packages

```
[1]: import pandas as pd
     import numpy as np
     import matplotlib.pyplot as plt
     import seaborn as sns
     import warnings
     import psycopg2
     import prophet
     from fbprophet import Prophet
     from fbprophet.plot import plot_plotly
     import plotly.offline as py
     import os
     warnings.filterwarnings('ignore')
     %matplotlib inline
     py.init_notebook_mode()
     color = sns.color_palette()
     pd.set_option('display.float_format', lambda x: '%.3f' % x)
     sns.set(rc={'figure.figsize':(8,5)})
```

0.0.2 Connecting to the Zapier database

- A view called **active_users_analysis** has been created in the **jchumley** schema using the *tasks_used_da table*.

- We'll use the *psycopg* odbc module to connect to the Redshift cluster. Pandas will be used extensively in conjunction with Postgres/Redshift queries.

- The user, password, host, and database parameters are loaded from a .bashrc file.

```
[2]: port='5439'
     user=os.environ['ZAP_USER']
     password=os.environ['ZAP_PASS']
     host=os.environ['ZAP_HOST']
     db=os.environ['ZAP_DB']
     conn = psycopg2.connect(user=user, password=password, host=host, port=port,␣
      ↪database=db)
```

0.0.3 Creating dataframe

- Let's create our first dataframe using pandas read_sql method.

```
[3]: query = """
SELECT * FROM
jchumley.active_users_analysis
"""
df = pd.read_sql(query, conn)
df.head()
```

```
[3]:         date  user_id  account_id  sum_tasks_used  cum_sum_tasks_used  \
0  2017-01-01        1           1              48                  48
1  2017-01-02        1           1              65                 113
2  2017-01-03        1           1              71                 184
3  2017-01-04        1           1              64                 304
4  2017-01-05        1           1              74                 406

   num_days_since_last_active  active  churned  month_num     month  dow_num  \
0                           0    True    False      1.000   January        0
1                           1    True    False      1.000   January        1
2                           1    True    False      1.000   January        2
3                           1    True    False      1.000   January        3
4                           1    True    False      1.000   January        4

         dow
0     Sunday
1     Monday
2    Tuesday
3  Wednesday
4   Thursday
```

- What is the shape of this dataset?

```
[4]: df.shape
```

```
[4]: (10547587, 12)
```

- There are ~10.5m rows and 12 columns.

- We'll also compute the ratio of missing values for each column.

```
[5]: df_na = (df.isnull().sum() / len(df)) * 100
df_na = df_na.drop(df_na[df_na == 0].index).sort_values(ascending=False)
missing_data = pd.DataFrame({'Missing Ratio' :df_na})
missing_data.head()
```

```
[5]: Empty DataFrame
Columns: [Missing Ratio]
```

```
Index: []
```

- We don't have null values in this data set.

0.1 Exploratory data analysis

0.1.1 Correlation Matrix

- Let's gain a deeper understanding of the associations between the features in this dataset. We'll use the pandas *corr* method to complete this task.

[6]: `df.corr()`

[6]:

	user_id	account_id	sum_tasks_used \
user_id	1.000	0.974	-0.022
account_id	0.974	1.000	-0.019
sum_tasks_used	-0.022	-0.019	1.000
cum_sum_tasks_used	0.984	0.958	-0.022
num_days_since_last_active	0.021	0.020	-0.012
active	-0.008	-0.008	0.004
churned	0.007	0.006	-0.003
month_num	0.108	0.108	0.001
dow_num	0.002	0.002	0.000

	cum_sum_tasks_used	num_days_since_last_active \
user_id	0.984	0.021
account_id	0.958	0.020
sum_tasks_used	-0.022	-0.012
cum_sum_tasks_used	1.000	0.023
num_days_since_last_active	0.023	1.000
active	-0.009	-0.760
churned	0.008	0.519
month_num	0.096	0.060
dow_num	0.002	-0.031

	active	churned	month_num	dow_num
user_id	-0.008	0.007	0.108	0.002
account_id	-0.008	0.006	0.108	0.002
sum_tasks_used	0.004	-0.003	0.001	0.000
cum_sum_tasks_used	-0.009	0.008	0.096	0.002
num_days_since_last_active	-0.760	0.519	0.060	-0.031
active	1.000	-0.873	-0.041	-0.000
churned	-0.873	1.000	0.029	0.000
month_num	-0.041	0.029	1.000	0.016
dow_num	-0.000	0.000	0.016	1.000

- For this stage of analysis, the user_id and account_id columns will be ignored since they are unique identifiers.

3

- Notice that there is a strong negative association between the features **active** and **churn** which makes sense since they are mutually exclusive.

0.1.2 Active and churned users by month

- Let's take a look at counts of active and churned users.
- We'll run the following aggregate query to view churn and active user counts for each month in the dataset.

```
[7]: query_month = """
SELECT
  month,
  month_num,
  count(CASE WHEN churned is TRUE then 1 END) as churned_count,
  count(CASE WHEN active is TRUE then 1 END) as active_count,
  ROUND(
    100.0 * count(CASE WHEN churned is TRUE then 1 END)/ count(*),
    2
  ) as churned_ratio,
  ROUND(
    100.0 * count(CASE WHEN active is TRUE then 1 END)/ count(*),
    2
  ) as active_ratio,
  count(*) as total_user_count
FROM
  jchumley.active_users_analysis
GROUP BY
  month,
  month_num
ORDER BY
  month_num
"""
df_month = pd.read_sql(query_month, conn)
```

```
[8]: df_month
```

```
[8]:        month  month_num  churned_count  active_count  churned_ratio  \
     0    January      1.000            120       1927467          0.010
     1   February      2.000           8256       1881695          0.440
     2      March      3.000          12974       2207687          0.580
     3      April      4.000          12676       2100295          0.600
     4        May      5.000          14207       2282399          0.620
     5       June      6.000            964         83651          1.140

         active_ratio  total_user_count
     0         99.990           1927587
     1         99.560           1889981
     2         99.300           2223331
     3         99.170           2117882
```
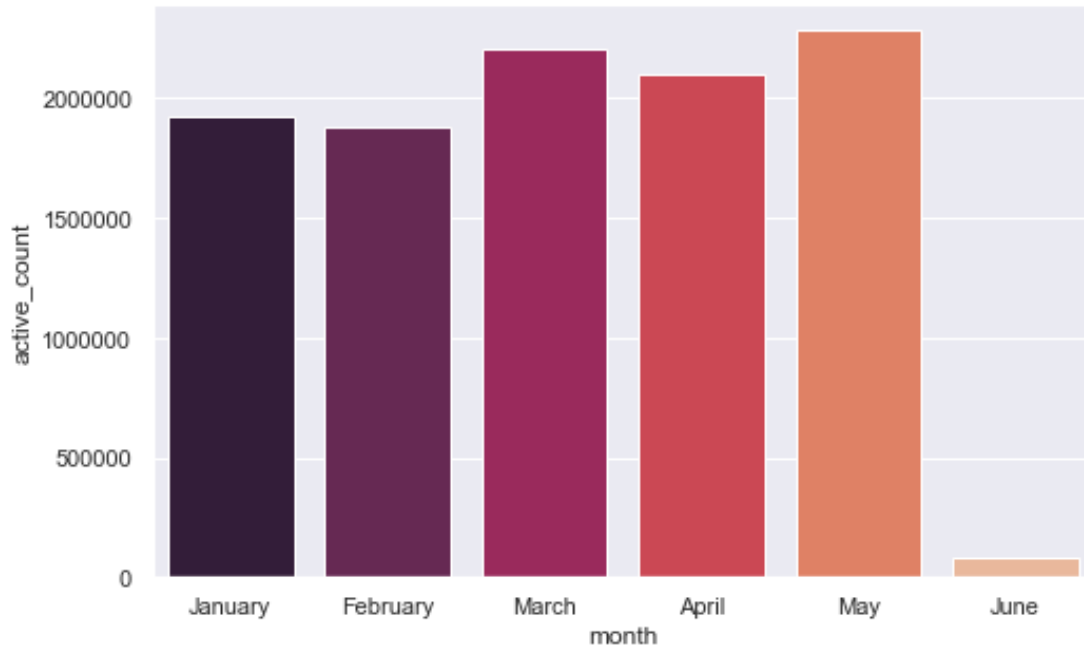
```
4        99.070          2303880
5        98.500            84926
```

[9]: `sns.barplot(x='month', y='active_count', palette="rocket", data=df_month)`

[9]: `<matplotlib.axes._subplots.AxesSubplot at 0x1a2d769150>`

- We can see that May had the highest count of active users. June has a very low count, perhaps because of data availability. Let's verify.

[10]:
```
query_months_date_counts = """
SELECT
  month,
  COUNT(
    DISTINCT(date)
  )
FROM
  jchumley.active_users_analysis
GROUP BY
  month
"""
pd.read_sql(query_months_date_counts, conn)
```

[10]:
```
       month  count
0    January     31
1       June      1
```

```
2      May       31
3     April      30
4  February      28
5     March      31
```

- We can see that the dataset includes records from the first day of June 2017 but not other days in that month/year.

0.1.3 Churn by month

```
[11]: sns.barplot(x='month', y='churned_count', palette="rocket", data=df_month)
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2eaa9090>
```


- January's churn count would likely increase if records from 2016 were availabile.
- Potential factors contributing to June's relatively high churn count:

    1. Users either chose to not renew or cancel their subscriptions on the first day of a given month.
    2. All June records are from the first day of the month.

0.1.4 Day of week analysis

- We'll run a query similar to the one used to compute aggregates by month.

```
[12]: dow_query = """
SELECT
  dow,
  dow_num,
  count(CASE WHEN churned is TRUE then 1 END) as churned_count,
  count(CASE WHEN active is TRUE then 1 END) as active_count,
  ROUND(
    100.0 * count(CASE WHEN churned is TRUE then 1 END)/ count(*),
    2
  ) as churned_ratio,
  ROUND(
    100.0 * count(CASE WHEN active is TRUE then 1 END)/ count(*),
    2
  ) as active_ratio,
  count(*) as total_user_count
FROM
  jchumley.active_users_analysis
GROUP BY
  dow,
  dow_num
ORDER BY
  dow_num
"""
df_dow = pd.read_sql(dow_query, conn)
```

```
[13]: df_dow
```

```
[13]:           dow  dow_num  churned_count  active_count  churned_ratio  \
      0      Sunday        0           3619       1102657          0.330
      1      Monday        1           8207       1657615          0.490
      2     Tuesday        2           8524       1703666          0.500
      3   Wednesday        3           9268       1708853          0.540
      4    Thursday        4           8773       1689305          0.520
      5      Friday        5           6805       1540253          0.440
      6    Saturday        6           4001       1080845          0.370

         active_ratio  total_user_count
      0        99.570           1107429
      1        99.370           1668200
      2        99.340           1714992
      3        99.300           1720901
      4        99.320           1700834
      5        99.420           1549246
      6        99.530           1085985
```

0.1.5 Counts of active users by day of week

```
[14]: sns.barplot(x='dow', y='active_count', palette="rocket", data=df_dow)
```

```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2f89eed0>
```


0.1.6 Counts of churned users by day of week

```
[15]: sns.barplot(x='dow', y='churned_count', palette="rocket", data=df_dow)
```

```
[15]: <matplotlib.axes._subplots.AxesSubplot at 0x1a301f52d0>
```

- Since Zapier is a B2B platform, its users are expected to be most active while they are in the office.
- Recall that **churned** and **active** have a very strong negative association (see the correlation matrix above), so it's expected that churn is higher during normal business hours.

0.1.7 Power Users

- Let's see who the top 10 Zapier users based on the sum_tasks_used column.
- We'll query the view total_sum_tasks_used which has the cumulative total of sum of tasks used for each user from the tasks_used_da table.

```
[16]: power_user_query = """WITH total_sum_tasks_used_CTE AS (
        SELECT
          user_id,
          SUM(sum_tasks_used) as total_sum_tasks_used
        FROM
          jchumley.active_users_analysis
        GROUP BY
          user_id
        ORDER BY
          user_id
      )
      SELECT
        user_id,
        total_sum_tasks_used,
```

```
    ROUND(
      100.0 * total_sum_tasks_used /(
        SELECT
          SUM(total_sum_tasks_used) as total_sum_tasks_used_all_users
        FROM
          total_sum_tasks_used_CTE
      ),
      2
    ) pct_all_sum_tasks_used,
    RANK() OVER(
      ORDER BY
        total_sum_tasks_used DESC
    ) as user_rank_by_sum_tasks_used
FROM
  total_sum_tasks_used_CTE
  LIMIT 10
"""
df_power_users = pd.read_sql(power_user_query, conn)
df_power_users
```

[16]:

	user_id	total_sum_tasks_used	pct_all_sum_tasks_used	\
0	541993	15878183	3.120	
1	19415	12574205	2.470	
2	558906	8282192	1.630	
3	645563	6153468	1.210	
4	103042	3872036	0.760	
5	1830742	3305466	0.650	
6	166842	3136333	0.620	
7	289586	2659815	0.520	
8	617995	2335859	0.460	
9	115417	2174002	0.430	

	user_rank_by_sum_tasks_used
0	1
1	2
2	3
3	4
4	5
5	6
6	7
7	8
8	9
9	10

- Would be interesting to dig deeper into metadata on these users (e.g, age, industry, gender, profession, etc.)

0.2 Timeseries forecasting

- We'll use Prophet to forecast active users and churn into the next 365 days. See https://facebook.github.io/prophet/ if you're not familiar with Prophet.
- Let's run another aggregate query against the active_users_analysis view, this time we'll group by date.

```python
[17]: user_aggregates_query = """
SELECT
  date as ds,
  count(CASE WHEN churned is TRUE then 1 END) as churned_count,
  count(CASE WHEN active is TRUE then 1 END) as active_count,
  ROUND(
    100.0 * count(CASE WHEN churned is TRUE then 1 END)/ count(*),
    2
  ) as churned_ratio,
  ROUND(
    100.0 * count(CASE WHEN active is TRUE then 1 END)/ count(*),
    2
  ) as active_ratio,
  count(*) as total_count
FROM
  jchumley.active_users_analysis
GROUP BY
  date
ORDER BY
  date"""
df_ts = pd.read_sql(user_aggregates_query, conn)
df_ts.head()
```

```
[17]:           ds  churned_count  active_count  churned_ratio  active_ratio  \
      0  2017-01-01              0         38903          0.000       100.000
      1  2017-01-02              0         53046          0.000       100.000
      2  2017-01-03              0         64350          0.000       100.000
      3  2017-01-04              0         65926          0.000       100.000
      4  2017-01-05              0         66304          0.000       100.000

         total_count
      0        38903
      1        53046
      2        64350
      3        65926
      4        66304
```

0.2.1 Forecast of Active Users

- Let's create a timeseries forecast of counts of active users into June 2018.

```
[18]: df_ts_active = df_ts.copy()
      df_ts_active[['ds', 'y']] = df_ts[['ds', 'active_count']]
```

Instantiating and Prophet object and making predictions into June 2018.

```
[19]: m = Prophet()
      m.fit(df_ts_active)
      active_users_future = m.make_future_dataframe(periods=365)
      active_users_future.tail()
```

```
INFO:fbprophet:Disabling yearly seasonality. Run prophet with
yearly_seasonality=True to override this.
INFO:fbprophet:Disabling daily seasonality. Run prophet with
daily_seasonality=True to override this.
```

```
[19]:           ds
      512 2018-05-28
      513 2018-05-29
      514 2018-05-30
      515 2018-05-31
      516 2018-06-01
```

```
[20]: active_users_forecast = m.predict(active_users_future)
      active_users_forecast[['ds', 'yhat', 'yhat_lower', 'yhat_upper']].tail()
```

```
[20]:           ds       yhat  yhat_lower  yhat_upper
      512 2018-05-28 93785.262   64103.243   124007.217
      513 2018-05-29 95813.956   66355.139   125734.392
      514 2018-05-30 95986.144   66377.062   126140.119
      515 2018-05-31 95036.977   65501.869   125028.040
      516 2018-06-01 91740.225   61635.807   122143.719
```

Visualizing our active users forecast

```
[21]: fig1_active = m.plot(active_users_forecast)
```

```
[22]: fig2 = m.plot_components(active_users_forecast)
```


13

- The model is predicting a steady growth of active users into June 2018.
- Notice that the weekly component aligns with our day of week analysis for active users.

0.2.2 Forecast of Churned Users

Similarly, let's create a timeseries forecast for churn into June 2018.

```
[23]: df_ts_churned = df_ts.copy()
      df_ts_churned[['ds', 'y']] = df_ts[['ds', 'churned_count']]
      df_ts_churned.head()
```

```
[23]:          ds  churned_count  active_count  churned_ratio  active_ratio  \
      0  2017-01-01              0         38903          0.000       100.000
      1  2017-01-02              0         53046          0.000       100.000
      2  2017-01-03              0         64350          0.000       100.000
      3  2017-01-04              0         65926          0.000       100.000
      4  2017-01-05              0         66304          0.000       100.000

         total_count  y
      0        38903  0
      1        53046  0
      2        64350  0
      3        65926  0
      4        66304  0
```

Instantiating and Prophet object and making predictions

```
[24]: m2=Prophet()
      m2.fit(df_ts_churned)
      churned_future = m2.make_future_dataframe(periods=365)
      churned_future.tail()
```

```
INFO:fbprophet:Disabling yearly seasonality. Run prophet with
yearly_seasonality=True to override this.
INFO:fbprophet:Disabling daily seasonality. Run prophet with
daily_seasonality=True to override this.
```

```
[24]:          ds
      512 2018-05-28
      513 2018-05-29
      514 2018-05-30
      515 2018-05-31
      516 2018-06-01
```

```
[25]: churned_user_forecast = m2.predict(churned_future)
      churned_user_forecast[['ds', 'yhat', 'yhat_lower', 'yhat_upper']].tail()
```

```
[25]:          ds      yhat  yhat_lower  yhat_upper
      512 2018-05-28   999.565     381.008    1624.277
      513 2018-05-29  1011.668     400.250    1662.279
```

```
514 2018-05-30 1043.238      464.941      1666.671
515 2018-05-31 1018.547      411.098      1668.066
516 2018-06-01  949.444      350.918      1561.290
```

Visualizing out churned users forecast

[28]: `fig1_churned = m2.plot(churned_user_forecast)`

[29]: `fig2_churned = m2.plot_components(churned_user_forecast)`

15

- Similar to the forecast of active users, churn is predicted to have steady growth into June 2018.
- Unlike the active user plot, we can see outliers in churn at the end of each month. This could be due to a large number of users who are registered under monthly billing subscriptions.

0.2.3 Business recommendations

- Identify new data sources containing Zapier users' metadata (e.g.,registration date, plan type, first usage date, acquisition channel, date of birth, gender, industry, etc.).
- Using additional data source combined with the task_used_da table, develop user metadata features.
- Develop a logistic regression model to generate churn probability scores for each user.
- Perform cohort analysis to group users by shared attributes.
- For users who have churned, send out a survey to understand why they have decided to no longer use the Zapier platform.
- Using the logistic regression model identify high-value power users who are at risks of churn.
- Run A/B testing experiments and compare churn rates between the test and control slices.
- Reach out to users who are on monthly subscription plans prior to the end of the month via email or another communication channel with the goal of customer retention.